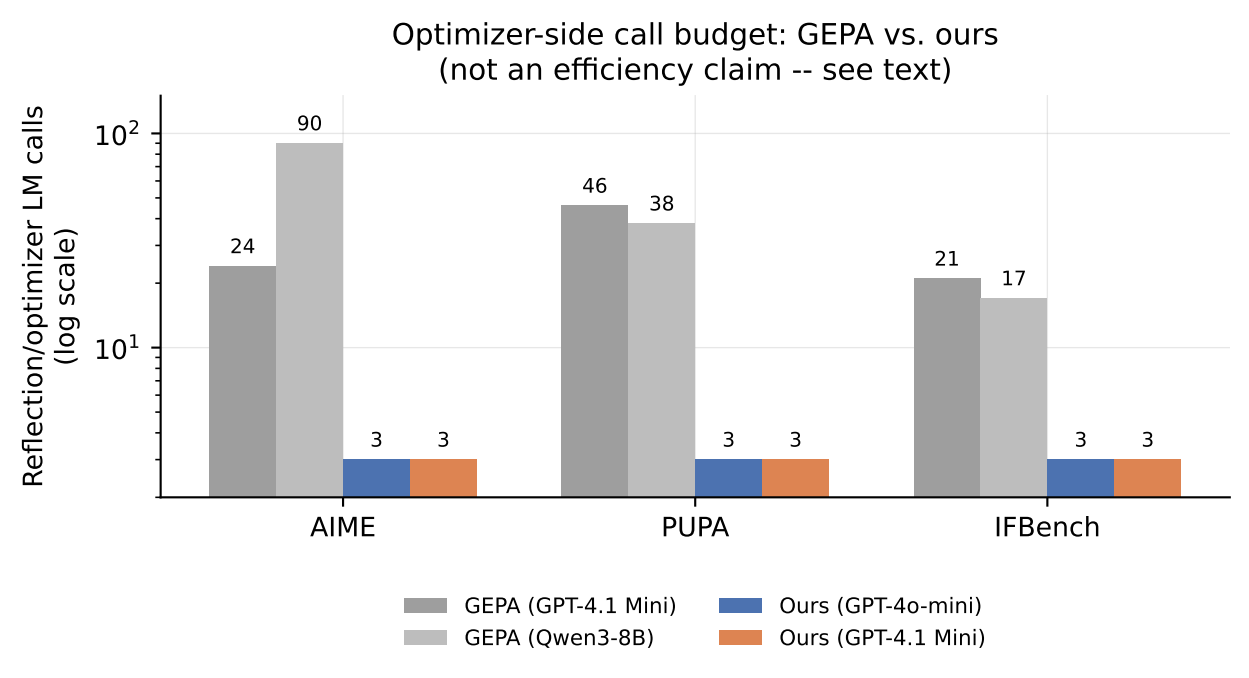
Optimizer-side call budget: GEPA vs. ours
(not an efficiency claim -- see text)
Reflection/optimizer LM calls
(log scale)
10^2
10^1
24
90
3
3
46
38
3
3
21
17
3
3
AIME
PUPA
IFBench
GEPA (GPT-4.1 Mini)
GEPA (Qwen3-8B)
Ours (GPT-4o-mini)
Ours (GPT-4.1 Mini)